

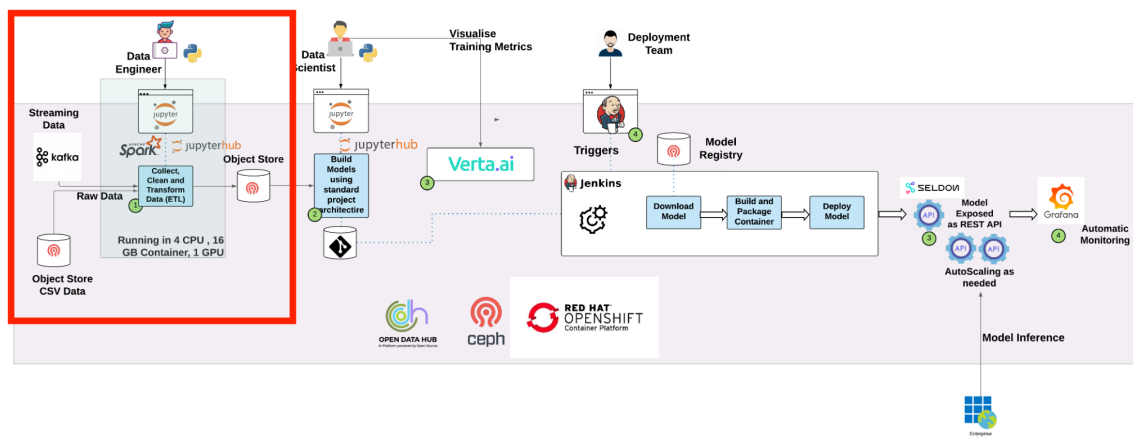
Lab 1 - Data Engineering

Introduction

In the previous lab, you explored the Open Data Hub component Superset - which provides easy to create charts and dashboards using the in-memory data engine Trino. Trino and Superset allow data residing anywhere to be accessed using a low latency in memory engine using SQL.

The Open Data Hub exposes a second data focused tool - for Extract Transform Load (ETL) of data originating in multiple data sources, i.e. Apache Spark. Spark allows finer grained ETL control than SQL/Trino does, e.g. using Regex to match data patterns. Spark provides a further toolset to allow data professionals to prepare quality data for consumption by data scientists and AI models.

This diagram illustrates the workflow we're implementing - the beginning part of the overall AI/ML workflow:

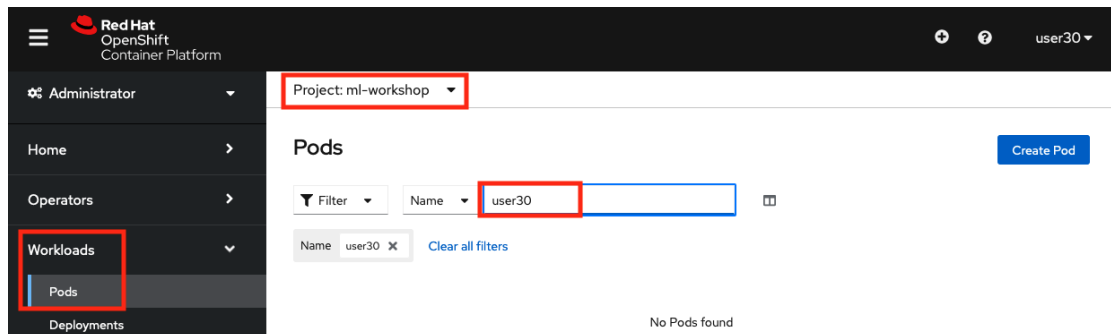


You can see, we source raw data from Kafka and S3 object storage. We use Jupyter notebooks to do some simple data engineering - combining these datasets on customerId using Spark. We then push that prepared data (a CSV file) to another bucket in our S3 object store, called Minio.

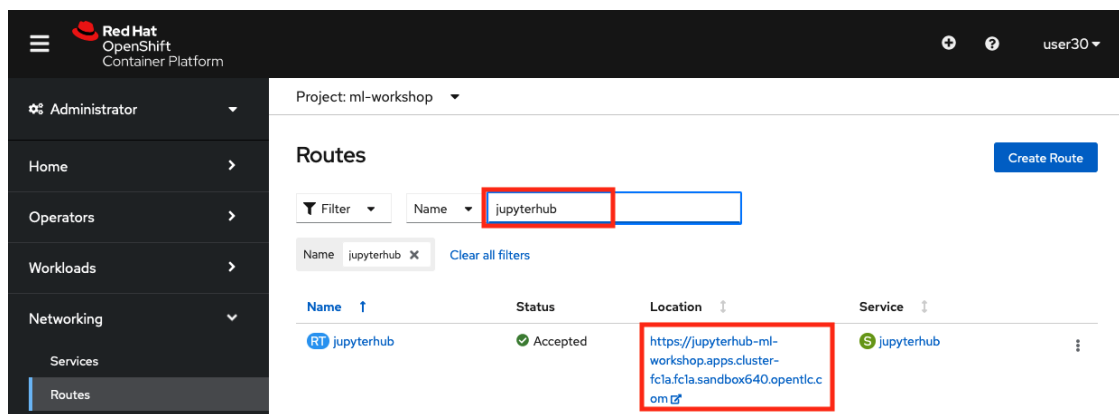
Instructions for the Spark workshop

Login to OpenShift using the credentials your administrator gave you. Open a second tab on your browser, also logged into OpenShift. Ensure your workshop project ml-workshop is selected.

1. In tab 1, Choose the Administration dropdown , navigate to Workloads -> Pods
Filter on your username, e.g. in my case user30. There won't be any pods shown - yet.



2. In tab 2, Choose the Administration dropdown , navigate to Networking -> Routes. Filter on *Jupyterhub* - and open the route.





This should open a new browser window requesting permission for jupyterhub-hub project to access your account. Select "Allow selected permissions".

Authorize Access

Service account jupyterhub-hub in project ml-workshop is requesting permission to access your account (user12

Requested permissions

☒ **user:info**

Read-only access to your user information (including username, identities, and group membership)

You will be redirected to https://jupyterhub-ml-workshop.apps.cluster-dhls2.dhls2.sandbox1573.opentlc.com/hub/oauth_callback

Allow selected permissions

Deny

. In a moment you'll see a screen similar to the following. Select

- *Minimal Python with Apache Spark* as the Notebook image.
- A *Large* container - allocating the maximum amount of CPU and memory available.

jupyterhub Home Token Services user30 Logout

Start a notebook server

Select options for your notebook server.

Notebook image

<input type="radio"/> MLWorkShop Notebook Image	<input type="radio"/> MLWorkShop Notebook Image Drift and Outlier
<input type="radio"/> Standard Data Science ? Python v3.8.3	<input type="radio"/> Elyra Notebook Image
<input type="radio"/> Minimal Python ? Python v3.8.3	<input type="radio"/> Minimal Python ? Python v3.6.8
<input type="radio"/> SciPy Notebook Image	<input checked="" type="radio"/> Minimal Python with Apache Spark
<input type="radio"/> Minimal Python with Apache Spark and SciPy	<input type="radio"/> Tensorflow Notebook Image

Deployment size

Container size

Large

Environment variables


[Add more variables](#)

Start server



Start Server. A few moments later the *files* view appears. The first thing we want to do is pull down our notebooks from our repository <https://github.com/masoodfaisal/ml-workshop>.

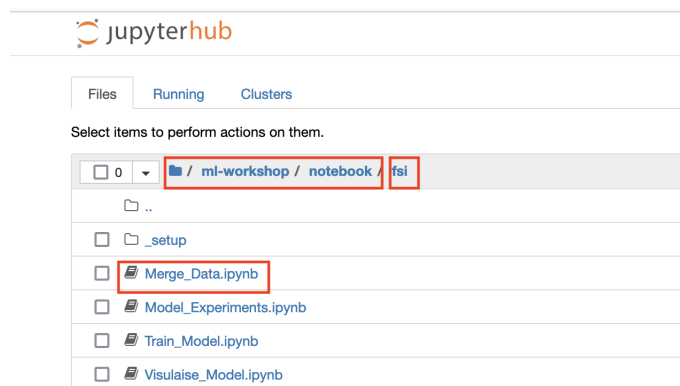
Choose New Terminal



Then paste this command into new Terminal and click enter:

```
nx,m/âz
```

Once done you can close that tab. Now refresh the *files* page - you'll see the *ml-workshop* folder. Now drill into ml-workshop -> notebook. Depending on your track **telco** or **fsi**, choose the appropriate subfolder (in this example **fsi**).



Now open up Merge_Data.ipynb, the file the data engineer uses to prepare their data.

Before we get going, you need to make a small change to the code.

Scroll down to the last cell and change the user to match the one provided by your instructor.

```
user_id = "<your username>"
```


Push prepared data to object storage and stop Spark cluster to save resources

Note - be sure to change this `user_id` on the next line to your username (something in the range user1 ... user30)

```
In [11]: user_id = "user29"
file_location = "s3a://data/full_data_csv" + user_id
dataFrom_All.repartition(1).write.mode("overwrite")\
.option("header", "true")\
.format("csv").save(file_location)
```

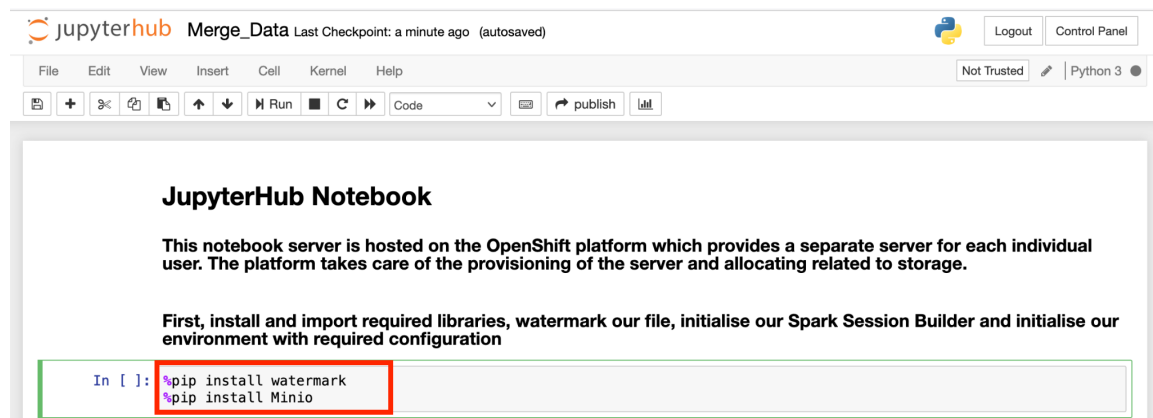
Change the `user_id` to be your `user_id` and save the notebook - using the save button at the top of the screen.

save the notebook - using the save button at the top of the screen:



You are now good to go!

Place the cursor inside the first cell with the `pip install` commands as shown.



To execute a cell, type the **SHIFT + RETURN** keys together. Walk through the entire file this way, executing as you go. (if there are any blank cells, skip through them using SHIFT + RETURN).

Here's a high level description of what's happening in the various cells:

```
1 In [ ]: %pip install watermark
          %pip install Minio
```

```
2 In [ ]: import os
          import json
          from pyspark import SparkConf
          from pyspark.sql import SparkSession, SQLContext
          from pyspark.sql.functions import from_json, col, to_json, struct
          import watermark
          from minio import Minio

          %matplotlib inline
          %load_ext watermark
```

```
3 In [ ]: %watermark -n -v -m -g -iv
```

```
In [ ]:
```

```
4 In [ ]: sparkSessionBuilder = SparkSession\
          .builder\
          .appName("Customer Churn ingest Pipeline")
```

```
5 In [ ]: os.environ['PYSPARK_SUBMIT_ARGS'] = \
          '--packages \
          org.postgresql:postgresql:42.2.10,\
          org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5,\
          org.apache.kafka:kafka-clients:2.4.0,\
          org.apache.spark:spark-streaming_2.11:2.4.5,\
          org.apache.hadoop:hadoop-aws:2.7.3 \
          --conf spark.jars.ivy=/tmp \
          --conf spark.hadoop.fs.s3a.endpoint=http://minio-ml-workshop:9000 \
          --conf spark.hadoop.fs.s3a.access.key=minio \
          --conf spark.hadoop.fs.s3a.secret.key=minio123 \
          --conf spark.hadoop.fs.s3a.path.style.access=true \
          --conf spark.hadoop.fs.s3a.impl=org.apache.hadoop.fs.s3a.S3AFileSystem \
          --master spark://' + os.environ['SPARK_CLUSTER'] + ':7077 pyspark-shell '
```

Connect to Spark Cluster provided by OpenShift Platform

```
6 In [ ]: spark = sparkSessionBuilder.getOrCreate()
          spark.sparkContext.setLogLevel("INFO")
          print('Spark context started.')
```

1. *pip install xxxx*, installs various libraries that aren't contained in our case container image
2. import the python libraries we need
3. *watermark* outputs the versions of various components, libraries, operating system attributes etc.
4. Here we create a Spark session, a precursor to firing up our own Spark server.
5. Here we set up various environment variables, including connection access to our S3 object store, in our case implemented using the open-source component Minio.
6. Here we actually start our Spark server. This cell can take several minutes to start.

```
7 In [ ]: dataframe_Customer = spark.read\
          .options(delimiter=',', inferSchema='True', header='True') \
          .csv("s3a://rawdata/Customer-Churn_P1.csv")
          dataframe_Customer.printSchema()

8 In [ ]: # dataframe_Products = spark.read\
          # .options(delimiter=',', inferSchema='True', header='True') \
          # .csv("s3a://rawdata/Customer-Churn_P2.csv")
          # dataframe_Products.printSchema()

9 In [ ]: from pyspark.sql.types import *
          from pyspark.sql.functions import *

          srcKafkaBrokers = "odh-message-bus-kafka-bootstrap:9092"
          srcKafkaTopic = "data"

          ..
          ..

10 In [ ]: dataFrom_All = dataframe_Customer.join(df0bj, "customerID", how="full")

          Push prepared data to object storage and stop Spark cluster to save resources

          Note - be sure to change this user_id on the next line to your username (something in the range user1 ...
          user30)

11 In [ ]: user_id = "user29"
          file_location = "s3a://data/full_data_csv" + user_id
          dataFrom_All.repartition(1).write.mode("overwrite")\
          .option("header", "true")\
          .format("csv").save(file_location)

12 In [ ]: spark.stop()
```

7. Here we pull in our data from S3 – our CSV based demographic data for each of our approximately 7000 customers.
8. Commented out – ignore
9. Here we pull in our data from Kafka – our product consumption data for each of our approximately 7000 customers.
10. We join these 2 datasets, on the common column to each: *customerID*.
11. We push our data to our object store – filename contains our username.
12. We are all done now – we stop our Spark server.

This is our data engineering workshop finished. It's a simple exercise, though the same toolset could be used for much more complex data engineering tasks.

Before we move on, as evidence of this self provisioned cluster, dedicated entirely to you as a user, move back to the Pods view, you saw above, keeping your username as a filter. Notice OpenShift has created a 3-node Spark cluster for us:

Administrator
Home
Operators
Workloads
Pods
Deployments
DeploymentConfigs
StatefulSets
Secrets
ConfigMaps
CronJobs
Jobs
DaemonSets
ReplicaSets
ReplicationControllers

Project: ml-workshop

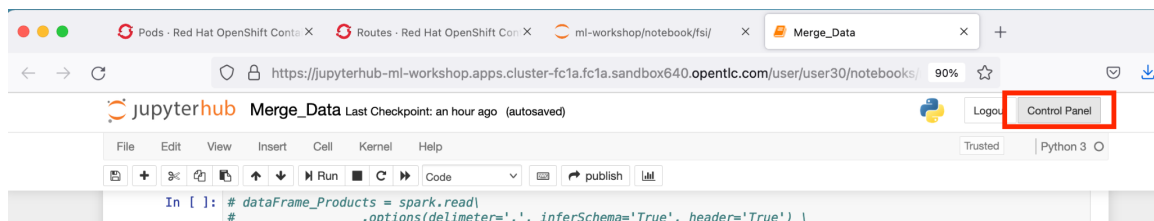
Pods

Filter
Name
user30

Name user30
Clear all filters

Name ↑	Status ↑	Ready ↑	Restarts ↑	Owner ↑	Memory ↑	CPU ↑
jupyterhub-nb-user30	Running	1/1	0	No owner	119.8 MiB	0.000 cores
spark-cluster-user30-m-qg6qk	Running	1/1	0	spark-cluster-user30-m	174.9 MiB	0.003 cores
spark-cluster-user30-w-65l98	Running	1/1	0	spark-cluster-user30-w	164.2 MiB	0.003 cores
spark-cluster-user30-w-l5vtd	Running	1/1	0	spark-cluster-user30-w	156.3 MiB	0.003 cores

Now we need to close down our Jupyter server. Choose *Control Panel* and shown - then on the next screen, choose Stop My Server.



The screenshot shows the JupyterLab interface with the URL `https://jupyterhub-ml-workshop.apps.cluster-fc1a.fc1a.sandbox640.opentlc.com/user/user30/notebooks/`. The **Control Panel** button is highlighted in the top right corner of the interface.

Immediately move back to your *Pods* screen - and observe your Spark pods being destroyed.

Name ↑	Status ↑	Ready ↑	Restarts ↑	Owner ↑	Memory ↑	CPU ↑
P spark-cluster-user30-m-qg6qk	Terminating	1/1	0	RC spark-cluster-user30-m	174.9 MiB	0.004 cores
P spark-cluster-user30-w-65l98	Terminating	1/1	0	RC spark-cluster-user30-w	164.6 MiB	0.004 cores
P spark-cluster-user30-w-l5vtd	Terminating	1/1	0	RC spark-cluster-user30-w	158.3 MiB	0.004 cores

This is a powerful demonstration of OpenShift's self service capabilities. No waiting for IT to provision you a server, no waiting around for access to a scarce Spark server. All self service, on demand, and those resources returned back to the central pool when finished.