

Lab 3 – ML OPs

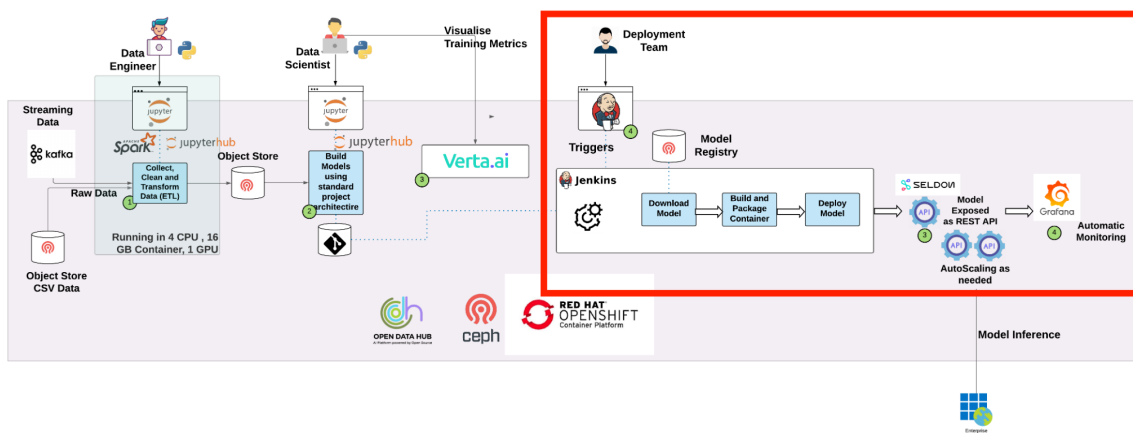
Introduction

Now as a data scientist, we've selected our chosen experiment, trained the model (a DecisionTreeClassifier) and finally we've pushed it to S3 Object storage.

Using the experiment id, as the identifier to retrieve from S3, we run our ML OPs pipeline and move the model and associated encoders to a *production* or *simulated production* environment. In this way, you add speed and potentially quality and security to this deployment process.

We also utilise the model serving component Seldon, which wraps the model behind a RESTful API, making the model easily available in this way for inference. As the flow is fully automated, we also ease and in fact eliminate the integration effort between application development and data scientists teams.

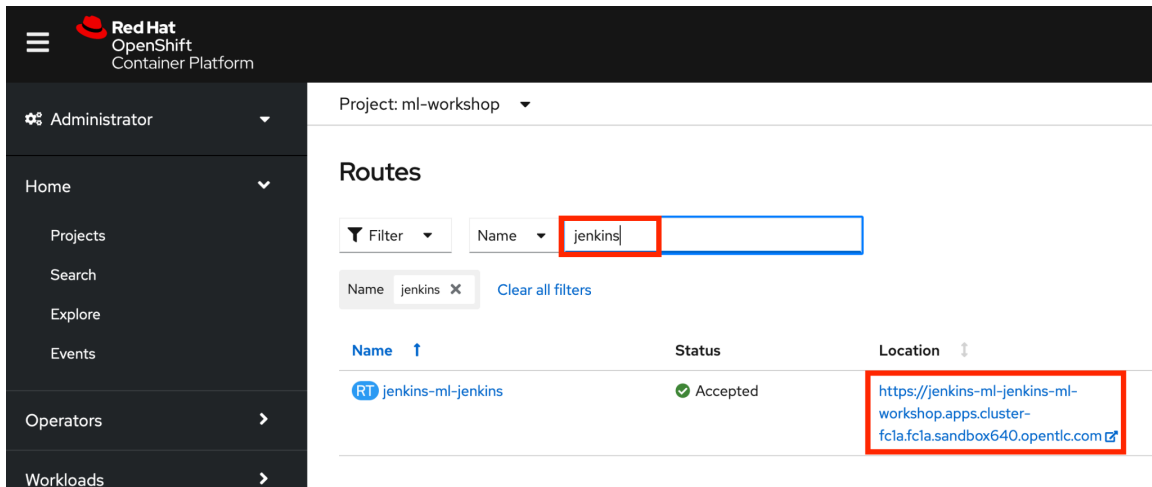
This diagram illustrates the workflow we're implementing – the ML OPs part of the overall AI/ML workflow:



Instructions to run the ML OPs workshop

Login to OpenShift using the credentials your administrator gave you. Ensure your workshop project ml-workshop is selected.

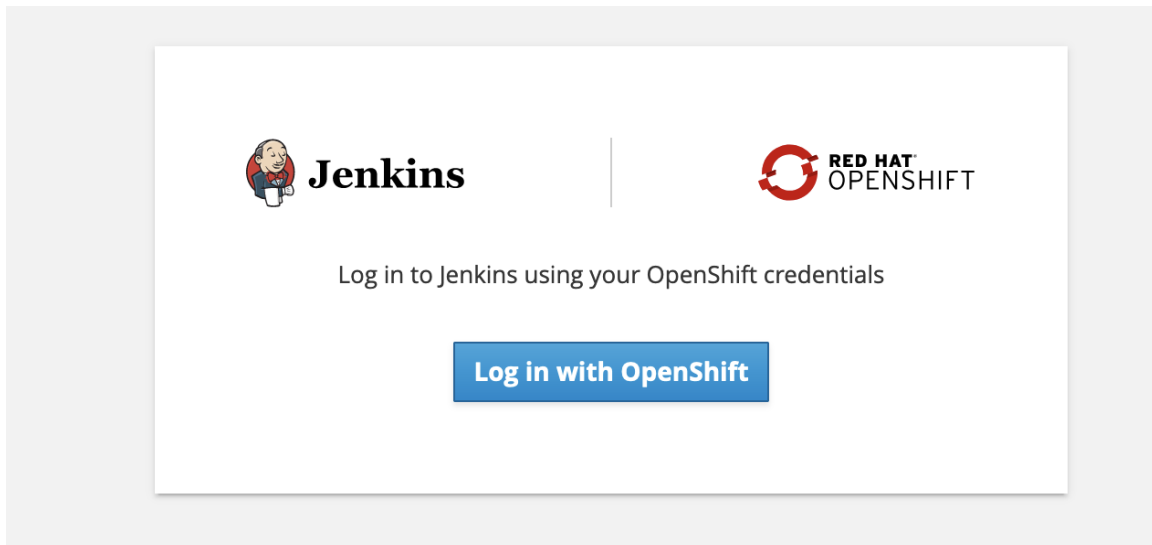
Choose the Administration dropdown , navigate to Network -> Routes. Filter on *jenkins* - and open the *Jenkins* route as shown.



The screenshot shows the OpenShift console interface. On the left is a sidebar with navigation options: Administrator, Home, Projects, Search, Explore, Events, Operators, and Workloads. The main area displays the 'Routes' page for the 'ml-workshop' project. A filter 'jenkins' is applied to the 'Name' column. A table lists the routes, with one entry 'jenkins-ml-jenkins' having a status of 'Accepted'. The 'Location' column for this route contains the URL 'https://jenkins-ml-jenkins-ml-workshop.apps.cluster-fcla.fcla.sandbox640.opentlc.com', which is highlighted with a red box.

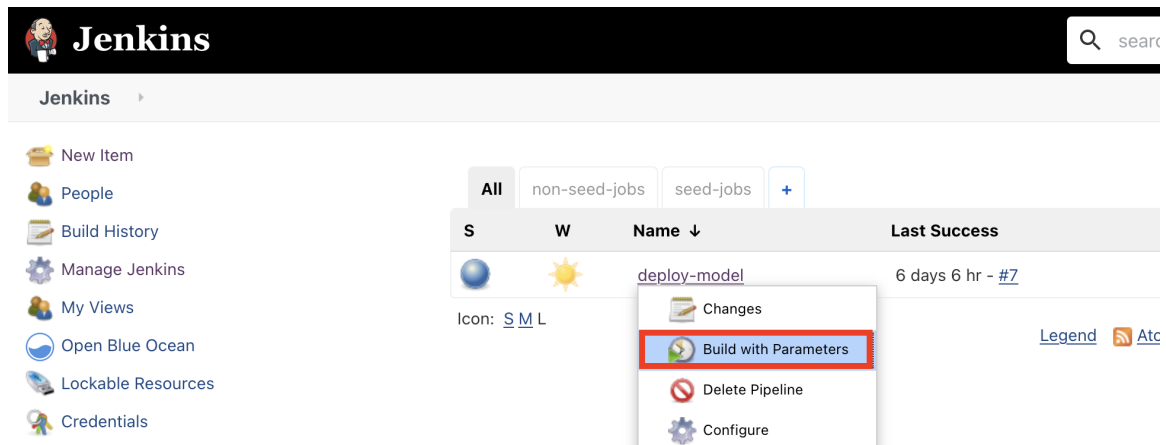
Name	Status	Location
jenkins-ml-jenkins	Accepted	https://jenkins-ml-jenkins-ml-workshop.apps.cluster-fcla.fcla.sandbox640.opentlc.com

When prompted enter your OpenShift credentials *userXX* and *openshift*, substituting *userXX* for your username.

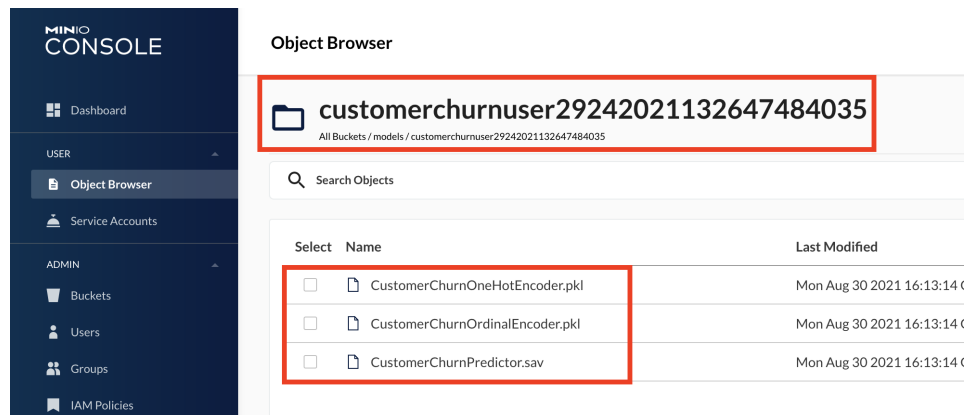


The screenshot shows the Jenkins login page. At the top, there are two logos: the Jenkins logo on the left and the Red Hat OpenShift logo on the right. Below the logos, the text 'Log in to Jenkins using your OpenShift credentials' is displayed. At the bottom, there is a blue button with the text 'Log in with OpenShift'.

Once in, you'll see a pipeline called `deploy-model`. Hover over it and an arrow will appear beside it. Choose **Build with Parameters** from the drop down menu.

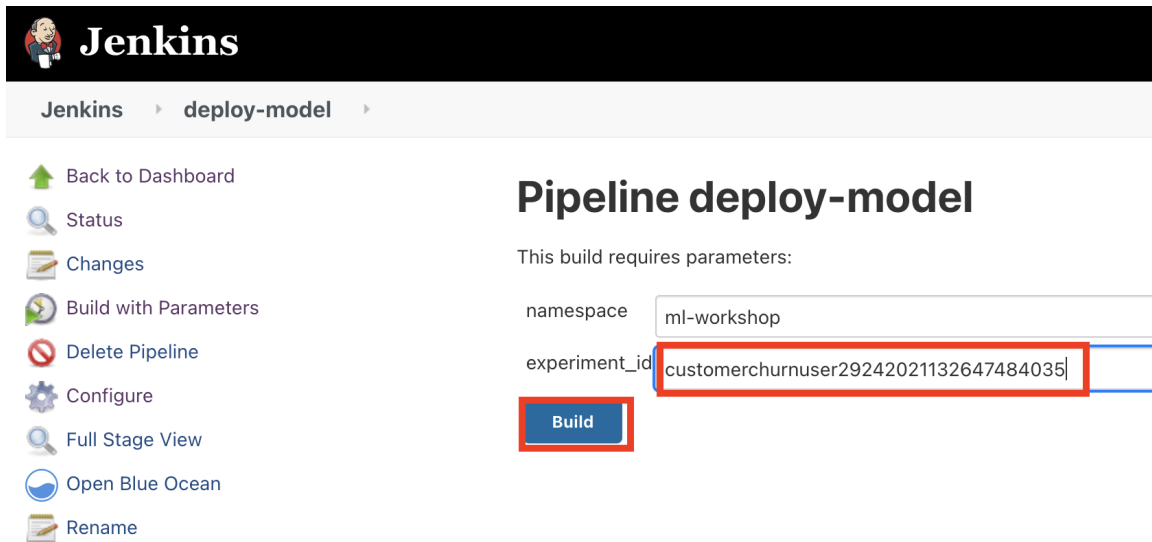


Remember the last exercise where you found your experiment outputs in the storage?



You need to use that experiment ID in this next step.

1. Enter the experiment id you retrieved at the end of the last workshop, in my case `customerchurnuser29242021132647484035`.
2. **Click Build**



Jenkins **deploy-model**

- Back to Dashboard
- Status
- Changes
- Build with Parameters
- Delete Pipeline
- Configure
- Full Stage View
- Open Blue Ocean
- Rename

Pipeline deploy-model

This build requires parameters:

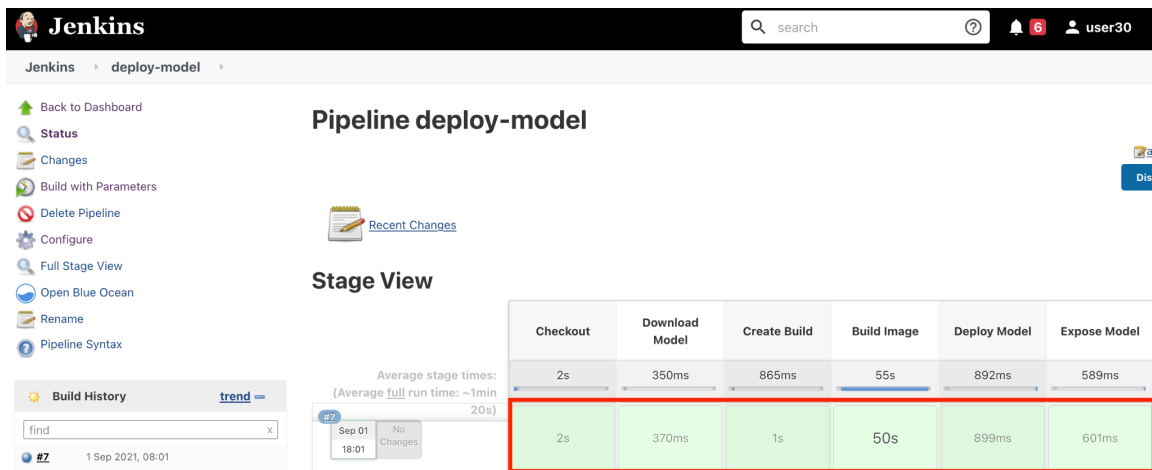
namespace

experiment_id

Build

Click **Build**

After a few seconds, you should see your pipeline begin and progress through the steps to completion you see highlighted here:



Jenkins **deploy-model**

Back to Dashboard

Status

Changes

Build with Parameters

Delete Pipeline

Configure

Full Stage View

Open Blue Ocean

Rename

Pipeline Syntax

Pipeline deploy-model

Recent Changes

Stage View

Average stage times:
(Average full run time: ~1min 20s)

Checkout	Download Model	Create Build	Build Image	Deploy Model	Expose Model
2s	350ms	865ms	55s	892ms	589ms
2s	370ms	1s	50s	899ms	601ms

Build History

find

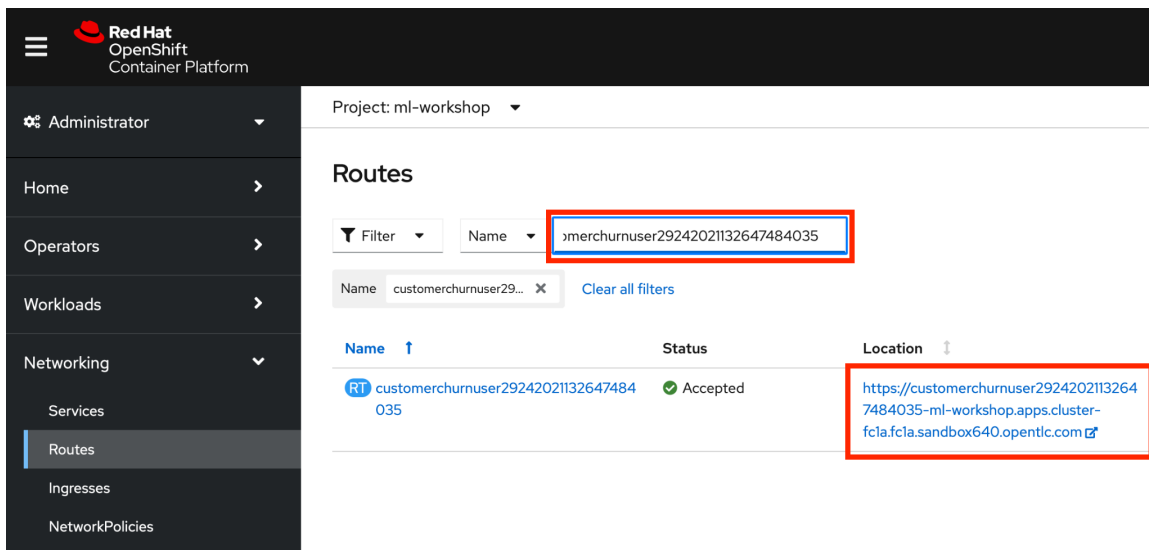
1 Sep 2021, 08:01

The pipeline then starts to package your model into an image and deploy it on to OpenShift. Proceed to the next section when the pipeline has completed.

Testing the Model via an API

You are now ready to test the API to your model.

1. Open the OpenShift tab on your browser.
2. Select the **Administrator** perspective in the left panel.
3. Click **Networking > Routes**
4. Filter on your experiment ID as shown below.



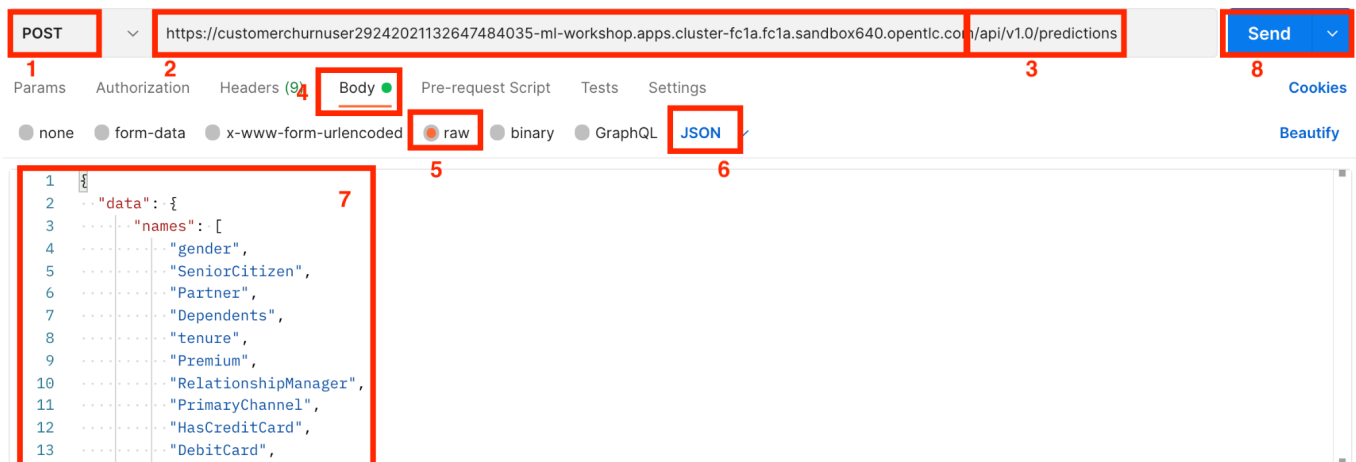
The screenshot shows the Red Hat OpenShift Container Platform Administrator interface. The left sidebar has a menu with 'Administrator' selected. The main content area is titled 'Routes' and shows a list of routes. A filter is applied to the 'Name' column, showing 'customerchurnuser29242021132647484035'. The first route in the list is highlighted, showing its name, status 'Accepted', and location URL: <https://customerchurnuser29242021132647484035-ml-workshop.apps.cluster-fc1a.fc1a.sandbox640.opentlc.com>.

The URL under the Locations column is the URL our pipeline has created for us - which we will use to make inference calls to our model. Make note of your URL, in my case

<https://customerchurnuser29242021132647484035-ml-workshop.apps.cluster-fc1a.fc1a.sandbox640.opentlc.com>

In order to make an inference call, you can use tools such as Postman, or the command line using curl, or there are various online options. I'll use *Postman* to illustrate.

Open Postman, and create a new Workspace and add a new Request to it. Populate it as follows:



1. Choose the POST method
2. Take your inference URL from the previous step
3. append the path **/api/v1.0/predictions** to that to form the full URL, in my case
`https://customerchurnuser29242021132647484035-ml-workshop.apps.cluster-fc1a.fc1a.sandbox640.opentlc.com/api/v1.0/predictions`
4. Select Body
5. Choose raw
6. and JSON as shown as the content type.
7. Paste the JSON located below in [Appendix 1 - Sample Inference Request Body](#) into the Body box.
Notice this represents a customer, we're asking the model to predict how likely it is they will churn. Notice also, we're passing in string values such as Brokerage etc.
Conversion using the *Ordinal* and *One-Hot* encoders to numeric values will be done by the running container exposing the API.
This simplifies these API calls greatly for Application developers making these inference calls.
8. Click Send to make your inference call

Here you see a sample response.

```
1  {
2    "data": {
3      "names": [
4        "t:0",
5        "t:1"
6      ],
7      "ndarray": [
8        [
9          0.8560606060606061,
10         0.14393939393939395
11        ]
12      ]
13    },
14    "meta": {
15      "model": "churn_model"
```

- There is an approximately 86% chance this customer will not churn (corresponding to t:0)
- There is an approximately 14% chance this customer will churn (corresponding to t:1)

Appendix 1 - Sample Inference Request Body

```
{
  "data": {
    "names": [
      "gender",
      "SeniorCitizen",
      "Partner",
      "Dependents",
      "tenure",
      "Premium",
      "RelationshipManager",
      "PrimaryChannel",
      "HasCreditCard",
      "DebitCard",
      "IncomeProtection",
      "WealthManagement",
      "HomeEquityLoans",
      "MoneyMarketAccount",
      "CreditRating",
      "PaperlessBilling",
      "AccountType",
      "MonthlyCharges",
      "TotalCharges"
    ],
    "ndarray": [
      [
        "Female",
        1,
        "Yes",
        "No",
        0,
        "Yes",
        "Yes",
        "Mobile",
        "Yes",
        "Yes",
        "No",
        "No",
        "Yes",
        "Yes",
        "Low",
        "No",
        "Brokerage",
        100,
        300
      ]
    ]
  }
}
```